


Lecture 3: Linear Filters




Signal Denoising

- Suppose we have a noisy 1D signal $f(x)$.
- For example, it could represent a company's stock price over time.
- In order to see the overall trend of the data, economists "smooth" out the noise using a moving average.
- A 3-point moving average will look at each data point and one point to either side.
- We then average the 3 original points and write this into a new signal $g(x)$.
- We then move to the next data point and repeat.






3-Point Moving Average



1D Mean Filter

- In signal processing, this **moving average** process is called a **1D mean filter**.
- We could do a 3-point mean filter where the weights are $w = \begin{bmatrix} \frac{1}{3} & \frac{1}{3} & \frac{1}{3} \end{bmatrix}$
- Or we could do a 5-point mean filter: $w = \begin{bmatrix} \frac{1}{5} & \frac{1}{5} & \frac{1}{5} & \frac{1}{5} & \frac{1}{5} \end{bmatrix}$
- In general, we can compute a **N-point filter** where N is an **odd** integer: $w = \begin{bmatrix} \frac{1}{N} & \frac{1}{N} & \frac{1}{N} & \dots & \frac{1}{N} \end{bmatrix}$
- As N gets **larger**, the data gets more **"smooth."**



1D Convolution

- We can think of doing a moving average with any set of weights in the vector w .
- We express this **moving average** procedure as a **convolution** denoted by $*$.

$$g(x) = (f * w)(x) = \sum_n f(n)w(x - n)$$


- A procedure that can be written as a convolution is called a linear filter.
- The weights w are called the filter or kernel.
- A 1D convolution can be computed using the Matlab command **conv**.

1D Gaussian Filter

- We could compute a weighted average by choosing any weights w that sum to 1.
- For example, we might want to give **more emphasis** to the **center** point and **less weight to far-away** points.
- The weights could be sampled from a **zero-mean Gaussian distribution** (bell curve).

$$G_{\sigma}(x) = \frac{1}{\sigma\sqrt{2\pi}} e^{-\frac{x^2}{2\sigma^2}}$$

- The **standard deviation** σ controls the **width** of the Gaussian.
- The **area** under the Gaussian is always **one**.



1D Gaussian Filter

- For Gaussian weights, the number of points N does not matter too much. The σ value controls the distribution.

$N=9, \sigma = 0.5$


$w =$ 0 0 0.0003 0.1065 0.7866 0.1065 0.0003 0 0

$N=9, \sigma = 2.0$

$w =$ 0.0276 0.0663 0.1238 0.1802 0.2042 0.1802 0.1238 0.0663 0.0276

center

- As $\sigma \rightarrow 0$, the Gaussian filter makes little change to the data. As σ gets larger, it resembles a mean filter.



2D Convolution

- A 2D convolution is a **weighted average** of a image neighborhood.
- Generally the neighborhood is a $N \times N$ square, where **N** is an **odd** integer.

$$g(x, y) = (f * w)(x, y) = \sum_m \sum_n f(m, n) w(x - m, y - n)$$

g = filtered image

f = original image

w = **filter** (**weights**)

2D Convolution

- Digitally, we slide the filter w around the image f and compute the weighted average of that neighborhood.
- To illustrate convolution with a 3x3 neighborhood, let's compute $g = f * w$ at pixel 13 below.

$W =$

w_1	w_2	w_3
w_4	w_5	w_6
w_7	w_8	w_9

$f =$

f_1	f_2	f_3	f_4	f_5
f_6	f_7	f_8	f_9	f_{10}
f_{11}	f_{12}	f_{13}	f_{14}	f_{15}
f_{16}	f_{17}	f_{18}	f_{19}	f_{20}
f_{21}	f_{22}	f_{23}	f_{24}	f_{25}

$$g_{13} = w_1 f_7 + w_2 f_8 + w_3 f_9 + w_4 f_{12} + w_5 f_{13} + w_6 f_{14} + w_7 f_{17} + w_8 f_{18} + w_9 f_{19}$$

- We then slide the 3x3 box to the next pixel and compute the weighted average of that neighborhood.

Boundary Conditions

- What happens when our neighborhood window partly off the side of the image?
- There are several choices for the boundary conditions (BCs).

Dirichlet BCs

0	0	0		
0	12	13	21	39
0	41	33	88	62
	32	44	71	55
	16	64	10	28

Periodic BCs

28	16	64		
39	12	13	21	39
62	41	33	88	62
	32	44	71	55
	16	64	10	28

Neumann (replicate) BCs

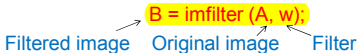
12	12	13		
12	12	13	21	39
41	41	33	88	62
	32	44	71	55
	16	64	10	28

Convolution in Matlab

- We compute a 2D convolution with the Matlab command **imfilter**.

B = imfilter (A, w);


Filtered image Original image Filter




- The **resulting** image B will have the **same size** as the image A.
- The **imfilter** command works on both **grayscale** and **color** images.
- The Matlab command **fspecial** can generate a variety of useful image filters.

Image Blur

- Applying a mean filter will "smooth" the edges of an image. We call this blur.



Mean Filter



Mean Filter

- In a $N \times N$ mean filter, all the weights are the same.
- As the window size N gets larger, the mean filter blurs the image more.




```
load woman;  
imagesc(X);
```

3x3 Mean Filter

```
w=1/9*ones(3,3);  
Y = imfilter(X,w);  
imagesc(Y);
```

5x5 Mean Filter

```
w=1/25*ones(5,5);  
Y = imfilter(X,w);  
imagesc(Y);
```



Gaussian Filter

- To build a Gaussian filter, we need to tell fspecial what **window size** to use and the **standard deviation** σ .
- To create a 5x5 Gaussian window with $\sigma = 0.8$:


```
G = fspecial('gaussian', [5,5], 0.8)
```

G =

0.0005	0.0050	0.0109	0.0050	0.0005
0.0050	0.0522	0.1141	0.0522	0.0050
0.0109	0.1141	0.2491	0.1141	0.0109
0.0050	0.0522	0.1141	0.0522	0.0050
0.0005	0.0050	0.0109	0.0050	0.0005

- It is sometimes helpful to display the filter weights.

- A Gaussian filter should be **high** in the **middle** (white) and low on the sides (black).




Gaussian Filter

- Although a Gaussian filter is good at removing Gaussian noise, it also blurs the image.

```
load woman;  
f = imnoise(uint8(X), 'gaussian', 0, 0.02);    %Add noise  
G = fspecial('gaussian', [5,5], 0.5);  
g = imfilter(f,G);
```

↖ ↗
window size σ

- We'll talk more about denoising next week.
- Increasing σ will increase the amount of blur.



Motion Blur

- We can simulate a moving object or camera by adding motion blur.
- We specify the angle indicating direction of motion (in degrees) and the number of pixels moved.

length of motion (in pixels)

angle of motion direction (in degrees)

`w = fspecial('motion', 6, 45)`

0	0	0	0.0438	0.1004
0	0	0.0438	0.1496	0.0438
0	0.0438	0.1496	0.0438	0
0.0438	0.1496	0.0438	0	0
0.1004	0.0438	0	0	0

Motion Blur



Laplacian Filter

- The Laplacian Filter has a large negative value in the center.
- The Laplacian approximates the second derivatives.

`w=fspecial('laplacian',0)`

0	1	0
1	-4	1
0	1	0

- Note these weights do not sum to 1.
- For any filter that contains negative values, we should change the image to double format.

`A = double(A);`

`B = imfilter(A, w);`

Laplacian Filter

- The Laplacian Filter detects edges, taking on negative value near edges.
- Subtracting the Laplacian filtered image from the original image can help sharpen edges and remove blur.

`w = fspecial('laplacian', 0);`

`A = double(A); B = A - imfilter(A, w);`

Original



Laplacian



Subtract Laplacian



Laplacian Filter

- If the image is **noisy**, the Laplacian filter will **pick up on** the noise.
- Subtracting the **Laplacian** will just **make the noise worse**.

$A = \text{imnoise}(A);$ % Add noise to image.


$w = \text{fspecial}('laplacian', 0);$

$A = \text{double}(A); \quad B = A - \text{imfilter}(A, w);$

Original



Laplacian



Subtract Laplacian



Prewitt Filter

- The Prewitt filter is designed to detect horizontal edges.
- Its transpose detects vertical edges.

`w = fspecial('prewitt')`

1	1	1
0	0	0
-1	-1	-1

`A = double(A);`

`H = imfilter(A,w);`

`V = imfilter(A,w');`



Difference of Gaussians (DoG)

- We can extract features of an image f by subtracting images blurred at two different levels.

$$G_{\sigma_1} * f - G_{\sigma_2} * f \quad \text{where } \sigma_1 < \sigma_2$$


Slightly blurry image

Very blurry image

```
G1 = fspecial('gaussian',[7,7],0.2);
```

```
G2 = fspecial('gaussian',[7,7],1.5);
```

```
DoG = imfilter(A,G1) - imfilter(A,G2);
```




Laplacian of Gaussian (LoG)

- A LoG filter is another popular way to locate image features.


`w = fspecial('log', [5,5], 0.8)`

0.0094	0.0470	0.0742	0.0470	0.0094
0.0470	0.0933	-0.0765	0.0933	0.0470
0.0742	-0.0765	-0.7770	-0.0765	0.0742
0.0470	0.0933	-0.0765	0.0933	0.0470
0.0094	0.0470	0.0742	0.0470	0.0094

Original



LoG



Unsharp Masking

- The curiously named Unsharp Filter enhances the edges in an image by subtracting off a Gaussian blurred image.

$$U * f = f - \alpha (G_{\sigma} * f)$$

- The unsharp filter is large in the middle and subtracts off the surrounding pixels.

`w = fspecial('unsharp', 0.5)`


-0.3333 -0.3333 -0.3333

-0.3333 3.6667 -0.3333

-0.3333 -0.3333 -0.3333

Unsharp Masking

Original



Small Alpha



Large Alpha



- Unsharp masking is a **deblurring** operation, but it is **not** really a **deconvolution** operation since it does not take into account the blur process.

The Convolution Theorem

- The Fourier Transform \mathcal{F} of a sequence f_n is defined as

$$\mathcal{F}(f)_k = \sum_{n=0}^{N-1} f_n e^{-\frac{2\pi i k n}{N}} \quad k=0, \dots, N-1$$

- A convolution is equivalent to pointwise multiplication of the Fourier transforms.

$$\mathcal{F}[f * w] = \mathcal{F}(f) \mathcal{F}(w)$$



- Then taking the Inverse Fourier Transform \mathcal{F}^{-1} gives

$$f * w = \mathcal{F}^{-1}[\mathcal{F}(f) \mathcal{F}(w)]$$

- A convolution on an image with N pixels can be computed in $O(N \log N)$ time.
- This means linear filters can be computed very quickly. This trick is built into the code of the `imfilter` command.

Point Spread Function (PSF)

- Every camera will have add some amount of blur.
- The Point Spread Function (PSF) is how the camera will respond to a point source.
- We think of the PSF as the blur kernel or filter weights inherent to that physical sensor.



Deconvolution

$$g = f * w$$

- Given g (observed image) and w (camera PSF), the goal of deconvolution is to recover the ideal image f .
- If we apply Fourier Transforms to both sides, the Convolution Theorem says:

$$\mathcal{F}(g) = \mathcal{F}[f * w] = \mathcal{F}(f) \mathcal{F}(w)$$

- Solve for f using the Inverse Fourier Transform \mathcal{F}^{-1}

$$\mathcal{F}(f) = \frac{\mathcal{F}(g)}{\mathcal{F}(w)} \Rightarrow f = \mathcal{F}^{-1} \left[\frac{\mathcal{F}(g)}{\mathcal{F}(w)} \right]$$

- This process is called Wiener Deconvolution.

Wiener Deconvolution

- The Matlab command `deconvwnr` performs deconvolution with a given PSF.

```
PSF = fspecial('gaussian', [5,5], 0.8);
```

```
blurred = imfilter(A, PSF);
```

```
deblurred = deconvwnr(blurred, PSF);
```

Create a Gaussian PSF.

Blur the image A with this PSF.
Then try to remove the blur.

Blurred Image



Wienered



Wiener Deconvolution

- Wiener Deconvolution is usually not practical because it is **extremely sensitive to noise**.

```
PSF = fspecial('gaussian', [5,5], 0.8);
```

```
blurred = imnoise(imfilter(A, PSF));
```


```
deblurred = deconvwnr(blurred, PSF);
```

Add small amount of noise to the blurred image.

Blurred Image + Noise



Wienered



Summary

- A linear filter is a **weighted average** around each pixel's neighborhood. We can write the filter as a **convolution**:

$$g = f * w.$$

- Linear filters have many uses:
 - Mean, Gaussian: Remove noise, but also blur the image.
 - Laplacian, Prewitt, DoG, LoG: Detect edges and other features.
 - Unsharp: Sharpens images. Can get similar result by subtracting a Laplacian filtered image.
- Wiener Deconvolution can remove blur, but only if the image was noise-free or if we know the signal-to-noise ratio (SNR).