


# Lecture 5: The Calculus of Images





## Discrete vs. Continuous

- Images are **discrete** objects. We are only given data at pixel locations.
- But many important geometry concepts are defined for continuous functions.
  - Derivatives and gradients
  - Area and volume
  - Curvature
  - Arc Length



# Images as Functions

- We can think of an image as a function of two variables  $f(x,y)$  defined on some rectangular domain  $\Omega$ .
- We know the value of the function at integer locations, e.g.  $f(2,3)$ .
- But what is the value at non-integer locations, like  $f(2.2, 3.4)$ ?
- Our data exists at discrete integer pixel locations. But we can *pretend* that the values exist in between the pixels.
- This allows us to discuss continuous concepts like derivatives and integrals on images.

# Discretization

- **Discretization** is the process of **approximating** a **mathematical concept** defined for **continuous** objects (like functions) into an **equivalent** concept for **discrete** objects (like images).

Continuous mathematical concept

Discretization

Discrete approximation

Example: **Riemann Sum**

$$\int_a^b f(x) dx$$



$$\sum_{i=1}^n f(x_i) \Delta x$$



# Finite Differences

- Let's start with a 1D signal  $f(x)$ .
- Recall the definition of the derivative.

$$f_x = \lim_{h \rightarrow 0} \frac{f(x+h) - f(x)}{h}$$

- But we can't let  $h$  go to zero. The smallest  $h$  can become is 1, because our data points are 1 pixel apart.
- So we *approximate* the derivative with  $h=1$ :  
$$f_x \approx f(x+1) - f(x)$$
- This type of approximation of the derivative is called a finite difference.

# Boundary Conditions

- So for our signal  $f(x)$ , we can approximate the derivative  $f_x$  at each point by looking at the difference with the next point.
- Suppose our vector has length  $n$ :  $n = \text{length}(f)$ ;
- What happens when we reach the last point?

$$f\_x(1) = f(2) - f(1);$$

$$f\_x(2) = f(3) - f(2);$$

$$f\_x(3) = f(4) - f(3);$$

....

$$f\_x(n-1) = f(n) - f(n-1)$$

$$f\_x(n) = ???$$

- Why does this code not work?  $f\_x = f(2:n) - f(1:n);$





# Boundary Conditions

- Neumann boundary conditions assumes an unknown derivative at the end of the data is zero.

$$f\_x(1) = f(2) - f(1);$$

$$f\_x(2) = f(3) - f(2);$$

$$f\_x(3) = f(4) - f(3);$$

....

$$f\_x(n) = f(n) - f(n) = 0;$$


- We can code this elegantly in one line of Matlab code using the colon operator. Just repeat the last entry.

$$f\_x = f([2:n,n]) - f(1:n);$$

OR

$$f\_x = f([2:n,n]) - f;$$





## Derivative of a Sine Wave

```
x = 0:0.1:2*pi;  
f = sin(x);  
n = length(f);  
f_x = f([2:n,n]) - f;  
subplot(121); plot(x, f);  
subplot(122); plot(x, f_x);
```



# Finite Difference Schemes

- There are several ways we could approximate the derivative  $f_x$ . Different approaches are called schemes.

- Forward Difference:**  $h=1$

$$f_x \approx D_x^+ f = f(x+1) - f(x)$$

$$f\_x = f([2:n,n]) - f;$$

- Backward Difference:**  $h=-1$

$$f_x \approx D_x^- f = f(x) - f(x-1)$$

$$f\_x = f - f([1,1:n-1]);$$

- Central Difference:**  $h=2$

$$f_x \approx D_x^0 f = \frac{f(x+1) - f(x-1)}{2}$$

$$f\_x = ( f([2:n,n]) - f([1,1:n-1]) ) / 2;$$

# Partial Derivatives

- An image is 2-dimensional, so we have a derivative in the **x-direction** and a derivative in the **y-direction**.
- Let  $f(x,y)$  be a **grayscale** image.
- The forward differences would give:

$[m,n] = \text{size}(f);$

$f\_x = f(:, [2:n, n]) - f;$

$f\_y = f([2:m, m], :) - f;$

$$f_x \approx D_x^+ f = f(x+1, y) - f(x, y)$$

$$f_y \approx D_y^+ f = f(x, y+1) - f(x, y)$$

# Partial Derivatives

- The derivative in **x-direction**  $u_x$  locates **vertical edges**.
- The derivative in **y-direction**  $u_y$  locates **horizontal edges**.

```
A = imread('cameraman.tif');  
A = double(A);  
[m,n] = size(A);  
A_x = A(:,[2:n,n]) - A;  
A_y = A([2:m,m],:) - A;  
subplot(121); imagesc(A_x);  
subplot(122); imagesc(A_y);
```

# Finite Differences as Filters

- You can think of a finite difference as a 3x3 linear filter applied to an image.

- Forward Difference:  $h=1$

$$f_x \approx D_x^+ f = f(x+1, y) - f(x, y)$$



- Backward Difference:  $h=-1$

$$f_x \approx D_x^- f = f(x, y) - f(x-1, y)$$



- Central Difference:  $h=2$

$$f_x \approx D_x^0 f = \frac{f(x+1, y) - f(x-1, y)}{2}$$





## Other Derivative Approximations

- The **Prewitt filter** uses **more pixels** so it is **less sensitive** to noise. But it **de-emphasizes values near the center**.

$$u_x \approx \begin{bmatrix} -1 & 0 & 1 \\ -1 & 0 & 1 \\ -1 & 0 & 1 \end{bmatrix} \quad u_y \approx \begin{bmatrix} 1 & 1 & 1 \\ 0 & 0 & 0 \\ -1 & -1 & -1 \end{bmatrix}$$

- The **Sobel filter** gives **more emphasis** to changes around the **center pixel**.

$$u_x \approx \begin{bmatrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{bmatrix} \quad u_y \approx \begin{bmatrix} 1 & 2 & 1 \\ 0 & 0 & 0 \\ -1 & -2 & -1 \end{bmatrix}$$


- There are many other finite difference schemes, like **upwind** and **minmod**. Each has pros and cons.

# The Gradient

- The gradient is a 2D vector listing the values of the partial derivatives at each point:

$$\nabla u = \langle u_x, u_y \rangle$$

- The gradient always points in the direction of **maximum positive change** (dark to light).




# Norm of Gradient

- The **norm** (magnitude) of the **gradient** vector tells us the total amount of change at each pixel.

$$\|\nabla u\| = \sqrt{u_x^2 + u_y^2}$$

- The norm of the gradient is large at edges of the image and zero in flat (single color) regions.



# Edge Detector

- We use the norm of the gradient to detect edges.

```
P = imread('pout.tif');  
P = double(P);  
[m,n] = size(P);  
Px = P(:,[2:n,n]) - P;  
Py = P([2:m,m],:) - P;  
N = sqrt(Px.^2 + Py.^2);  
imagesc(N);
```

Note the .^ for pointwise exponents.





## Second Derivatives

- To approximate a second derivative, we take a finite difference of a finite difference.

$$u_{xx} \approx D_x^-(D_x^+ u)$$

- Note we use one forward and one backward difference.

## 3 Ways to Code $u_{xx}$

- 1.) Forward then Backward Difference.

$$u_{xx} \approx D_x^-(D_x^+ u)$$

Dplus = u(:, [2:n, n]) - u;

u\_xx = Dplus - Dplus(:, [1, 1:n-1]);

- 2.) Backward then Forward Difference.

$$u_{xx} \approx D_x^+(D_x^- u)$$

Dminus = u - u(:, [1, 1:n-1]);

u\_xx = Dminus(:, [2:n, n]) - Dminus;

- 3.) Write out the formula.

$$u_{xx} \approx u(x+1, y) - 2u(x, y) + u(x-1, y)$$

u\_xx = u(:, [2:n, n]) - 2\*u + u(:, [1, 1:n-1]);

## Second Derivatives

```
[m,n] = size(u);  
% Second derivative in x: u_xx  
u_xx = u(:,2:n,n) - 2*u + u(:,1,1:n-1);  
% Second derivative in y: u_yy  
u_yy = u([2:m,m],:) - 2*u + u([1,1:m-1],:);  
% Diagonal derivative u_xy  
u_xy = ( u([2:m,m],[2:n,n])  
        + u([1,1:m-1],[1,1:n-1])  
        - u([1,1:m-1],[2:n,n])  
        - u([2:m,m],[1,1:n-1]) ) / 4;
```



# The Laplacian

- The Laplacian is the sum of the second derivatives:  
 $\Delta u = u_{xx} + u_{yy}$
- Recall we implemented a Laplacian filter.
- We use the Laplacian to locate edges. Subtracting the Laplacian sharpens the images.

$$\begin{array}{c} u_{xx} \\ \begin{array}{|c|c|c|} \hline & & \\ \hline & & \\ \hline & & \\ \hline \end{array} \end{array} + \begin{array}{c} u_{yy} \\ \begin{array}{|c|c|c|} \hline & & \\ \hline & & \\ \hline & & \\ \hline \end{array} \end{array} = \begin{array}{c} \Delta u \\ \begin{array}{|c|c|c|} \hline & & \\ \hline & & \\ \hline & & \\ \hline \end{array} \end{array}$$



# Double Integrals

- The double integral of  $u(x,y)$  over the domain  $\Omega$  is

$$\iint_{\Omega} u(x,y) \, dx \, dy$$

- The discrete approximation is simply a **double summation of all values** of  $u(x,y)$ :  
 **$d = \text{sum}(\text{sum}(u));$**
- Sometimes we get lazy and **vectorize** the variables as  **$\vec{x} = (x,y)$**  so we can write a single integral:

$$\int_{\Omega} u(\vec{x}) \, d\vec{x}$$

- But don't let this fool you, it's still a double sum!



# Measuring Noise

- We measured the noise levels last week using **SNR** and **RMSE**.
- But these statistics require an ideal noise-free image, which in general we don't have.
- We would like a way to judge how much noise an image contains **without requiring a magical perfect image.**

# Total Variation

- The Total Variation (TV) energy of an image  $u(x,y)$  is found by adding up the norm of the gradient (*Rudin-Osher-Fatemi, 1989*).

$$TV(u) = \iint_{\Omega} \|\nabla u\| \, dx \, dy$$

- We interpret TV as the total amount of jumps (variation) in the image.
- Or if we vectorize  $\vec{x} = (x, y)$ , we can write as

$$TV(u) = \int_{\Omega} \|\nabla u\| \, d\vec{x}$$

# 1D TV

- The 1-dimensional version for a function  $f(x)$  on  $[a,b]$  is


$$TV = \int_a^b |f'(x)| dx$$

- Ex Calculate the TV value of a sine wave on  $[0, 2\pi]$ .




## 2D TV

- Calculate TV energy for the image below.
- Assume the image is 100x200 pixels and the gray square is 30x30 pixels.




# The Co-Area Formula

- Co-Area Formula: The TV norm is equal to the perimeter of each shape times the jump at the perimeter.
- Suppose we have a circle of radius 5 pixels on a dark gray background. Calculate the TV energy value.



## Noise on TV

- Again suppose we have a circle of radius 5 pixels on a dark gray background.
- Let's add one noise pixel with a value 250.



# Measuring Noise

- TV does not tell us exactly how much noise is in the image, but if we have a version of an image with a high TV value then it is probably noisy.

TV Energy = 11,150,000



TV Energy = 1,830,000





## Your Very Own TV

- Ex Write a function that calculates the TV energy value of a grayscale image.

$$TV(u) = \iint_{\Omega} \|\nabla u\| \, dx \, dy$$

- We'll have to discretize this energy, so really we'll be computing an approximation of the TV energy.

# Curvature

- The curvature of a surface  $u(x,y)$  measures how quickly the unit tangent vector to the surface is changing:

$$\kappa = \nabla \cdot \left( \frac{\nabla u}{\|\nabla u\|} \right)$$

- Ex Write a function that computes the curvature matrix of a grayscale image. Watch out for division by zero!